
Escaping Groundhog Day

James MacGlashan

Department of Computer Science
Brown University
Providence, RI 02915
james_macglashan@brown.edu

Stefanie Tellex

Department of Computer Science
Brown University
Providence, RI 02915
stefie10@cs.brown.edu

Michael Littman

Department of Computer Science
Brown University
Providence, RI 02915
mlittman@cs.brown.edu

Abstract

The dominant approaches to reinforcement learning rely on a fixed state-action space and reward function that the agent is trying to maximize. During training, the agent is repeatedly reset to a predefined initial state or set of initial states. For example, in the classic RL Mountain Car domain, the agent starts at some point in the valley, continues until it reaches the top of the valley and then resets to somewhere else in the same valley. Learning in this regime is akin to the learning problem faced by Bill Murray in the 1993 movie *Groundhog Day* in which he repeatedly relives the same day, until he discovers the optimal policy and escapes to the next day. In a more realistic formulation for an RL agent, every day is a new day that may have similarities to the previous day, but the agent never encounters the same state twice. This formulation is a natural fit for robotics problems in which a robot is placed in a room in which it has never previously been, but has seen similar rooms with similar objects in the past. We formalize this problem as optimizing a learning or planning algorithm for a set of environments drawn from a distribution and present two sets of results for learning under these settings. First, we present *goal-based action priors* for learning how to accelerate planning in environments drawn from the distribution from a training set of environments drawn from the same distribution. Second, we present *sample-optimized Rademacher complexity*, which is a formal mechanism for assessing the risk in choosing a learning algorithm tuned on a training set drawn from the distribution for use on the entire distribution.

Keywords: Meta-learning, transfer learning, learning to plan,



Figure 1: Two different versions of the gold smelting task. The task on the right is only able to be solved after learning in the simpler version of the left because the state space is too large.

1 Introduction

The dominant approaches to reinforcement learning rely on a fixed state-action space and reward function that the agent is trying to maximize. During training, the agent is repeatedly reset to a predefined initial state or set of initial states. For example, in the classic RL Mountain Car domain, the agent starts at some point in the valley, continues until it reaches the top of the valley and then resets to somewhere else in the same valley. Learning in this regime is akin to the learning problem faced by Bill Murray in the 1993 movie *Groundhog Day* in which he repeatedly relives the same day, until he discovers the optimal policy and escapes to the next day. In a more realistic formulation for an RL agent, every day is a new day that may have similarities to the previous day, but the agent never encounters the same state twice. This formulation is a natural fit for robotics problems in which a robot is placed in a room in which it has never previously been, but has seen similar rooms with similar objects in the past. We formalize this problem as optimizing a learning or planning algorithm for a set of environments drawn from a distribution.

In both cases, we assume environments are defined by a Markov decision process (MDP). An MDP is defined by the tuple (S, A, T, R, s_0) , where S is the state space; A is the action set; $T(s'|s, a)$ is the transition dynamics, which specifies the probability of transitioning to state s' after taking action a in state s ; $R(s, a, s')$ is the reward function, which specifies the reward received by the agent for taking action a in state s and then transitioning to state s' ; and s_0 is an initial state.

The goal of planning or learning in an MDP is to find a (near-)optimal policy π that maps states to actions. A policy is typically considered optimal if following it from the initial state maximizes the expected discounted future reward. For a planning, the agent has complete access to the MDP and can search for a policy before acting in the world. In reinforcement learning (RL), the agent does not have access to the transition dynamics or reward function and must learn how to act through interaction.

In our learning setting, an agent is given a set of sample training MDPs on which to optimize its performance. In the case of planning, learning from the training set is used to decrease computation time. In RL, learning from the training set is used to decrease learning time in new environments.

We present results for the learning to plan and learning to learn setting. In the learning to plan setting we present learnable *goal-based action priors* that accelerate planning on related MDPs. In the learning to learn setting, we present *sample-optimized Rademacher complexity*, which is a formal mechanism for assessing the risk in choosing a learning algorithm tuned on a training set drawn from the distribution for use on the entire distribution.

2 Learning to Plan

Robots operating in unstructured, stochastic environments such as a factory floor or a kitchen face a difficult planning problem due to the large state space and the very large set of possible tasks [3, 6]. A powerful and flexible robot such as a mobile manipulator in the home has a very large set of possible actions, any of which may be relevant depending on the current goal (for example, robots assembling furniture [6] or baking cookies [3].) When a robot is manipulating objects in an environment, an object can be placed anywhere in a large set of locations. The size of the state space increases exponentially with the number of objects, which bounds the placement problems that the robot is able to expediently solve. Depending on the reward function (which is unknown before runtime), any of these states and actions may be relevant to the solution, but for any specific reward function, most of them are irrelevant. For instance, when making brownies, the oven and flour are important, while the soy sauce and sauté pan are not. For a different task, such as stir-frying broccoli, the robot must use a different set of objects and actions.

To confront this state-action space explosion, prior work has explored adding knowledge to the planner, such as options [10] and macro-actions [4, 8]. However, while these methods can allow the agent to search more deeply in the state space, they add non-primitive actions to the planner which *increase* the branching factor of the state-action space. The resulting augmented space is even larger, which can have the paradoxical effect of increasing the search time for a good policy [5].




Figure 2: Average results from all maps.

To address these issues, we learn action priors that are conditioned on the current state and goal from training MDPs drawn from a distribution. Because we condition on both the state and goal description, we refer to this goal-based action prior as a knowledge base of *affordances*. Affordances enable the robot to prune irrelevant actions from the search space on a state-by-state basis based on the agent’s current goal and focus on the most promising parts of the state space. Actions are determined to be irrelevant when according to the affordance model they have a low probability of being optimal.

To learn affordances, we sample a set of training worlds from the domain (W), for which the optimal policy, π , may be tractably computed using existing planning methods. Then, we compute the maximum likelihood estimate of the parameter vector θ_i for each action using the policy. In our experiments, we use a Naive Bayes model in which the probability of state features are treated as independent given the optimal action. During the learning phase, the agent learns which actions are useful under different conditions. At test time, the agent will see different, randomly generated worlds from the same domain, and use the learned affordances to increase its speed at inferring a plan. For simplicity, our learning process uses a strict separation between training and test; after learning is complete our model parameters remain fixed for the entire test time.

We evaluate our approach using the game Minecraft. Minecraft is a 3-D blocks game in which the user can place, craft, and destroy blocks of different types. Minecraft’s physics and action space allow users to create complex systems, including logic gates and functional scientific graphing calculators¹. Minecraft serves as a model for robotic tasks such as cooking assistance, assembling items in a factory, object retrieval, and complex terrain traversal. Our experiments consisted of five common tasks in Minecraft, including constructing bridges over trenches, smelting gold, tunneling through walls, basic path planning, and digging to find an object. Figure 1 shows two scenes from Minecraft that were drawn from our distribution of MDPs. We tested on randomized worlds of varying size and difficulty. The generated test worlds varied in size from tens of thousands of states to hundreds of thousands of states. The agent learned affordances from a training set consisting of 25 simple state spaces of each map type (100 total maps), each approximately a 1,000-10,000 state world. We conducted all tests with a single knowledge base. Learning this knowledge base took approximately one hour run in parallel on a computing grid.

We use Real-Time Dynamic Programming (RTDP) [1] as our baseline planner, a sampling-based algorithm that does not require the planner to visit all states. We compare RTDP with learned affordance-aware RTDP (LA-RTDP), and expert-defined affordance-aware RTDP (EA-RTDP), in which the action priors are specified by an expert. We terminated each planner when the maximum change in the value function was less than 0.01 for 100 consecutive policy rollouts, or the planner failed to converge after 1000 rollouts. The reward function was -1 for all transitions, except transitions to states in which the agent was in lava, where we set the reward to -10 . The goal was set to be terminal and the discount factor was $\gamma = 0.99$. To introduce non-determinism into our problem, movement actions (move, rotate, jump) in all experiments had a small probability (0.05) of incorrectly applying a different movement action. This noise factor approximates noise faced by a physical robot that attempts to execute actions in a real-world domain and can affect the optimal policy due to the existence of lava pits that the agent can fall into. Figure 2 summarizes the testing performance for all planning algorithms in terms of average number of Bellman updates in RTDP, the average cost of the policy found after planning termination, and the average amount of CPU time. In all cases, planning with learned affordances performed best, then expert performance, and finally the baseline performed worst. Although RTDP converge to the optimal policy in the limit, its average cost was worse than the other baselines because in many cases it was not able to find a reasonable policy in the amount of planning time allotted.

3 Learning to Learn

For the case of optimizing a learning algorithm for a distribution of MDPs, we seek the answer the question: How do you know if your reinforcement-learning (RL) algorithm is the right one for your problem? How can you tell if you are at risk for overfitting? Or underfitting?

¹<https://www.youtube.com/watch?v=wgJfVRhotlQ>




Figure 3: Training and testing performance for two distributions in the 5-state chain environment.

A core concept used in mitigating the effects of overfitting is to apply more powerful hypothesis classes only when copious data is available to fit their free parameters, restricting learners to weaker hypothesis classes when data is more sparse. In some well-studied cases, bounds relating the hypothesis class to the amount of data required to fit it accurately are known [2].

As an example, consider the 5-state chain [9]: a well known Markov decision process (MDP). This problem consists of a set of 5 states arranged in a linear order. Action a_1 causes a transition to the next state in the ordering. Action a_2 resets the state to the first in the chain. Action a_1 from the last state in the chain results in remaining in the state with a high reward (+10) and zero reward otherwise. Action a_2 always has a reward of +2. With probability 0.2 (the slip probability), however, the selected action has the effect of the non-selected action. Though tiny, this MDP is challenging for some learning algorithms because action a_2 acts as a temptation that keeps the agent from discovering the optimal policy (always take action a_1). We measure the performance of a learning algorithm in the environment by its probability of finding an optimal policy after 1000 steps of experience.

Consider two different variations of the Q-learning algorithm [11] that we might want to apply to this problem. In both, the learning rate (α) is set to a value between 0.0 and 0.5 and the exploration rate (ϵ) is set to a value between 0.0 and 0.4. In the *constant initialization* algorithm, all Q values are initially set to 45 (something in the vicinity of the likely final value function). In the *variable initialization* algorithm, each state-action pair is initialized independently to some value between 0 and 200. Note that the variable initialization algorithm subsumes the constant initialization algorithm since it can be configured to initially set all Q values to 45.

If we tune parameters to the 5-state chain MDP, the variable initialization algorithm will perform better. However, that does not mean the tuned variable initialization algorithm will perform better on a distribution of varying 5-state chain problems, depending on the properties of the distribution. For example, consider two different MDP distributions. In distribution 1, all MDPs are 5-state chains with states ordered consistently but slip probability varying between 0.19 and 0.21. In distribution 2, all MDPs are 5-state chains, with the order of states in the chain varying from MDP to MDP and slip probability varying between 0.00 and 0.50. By exhaustive testing shown in Figure 3, we find that for distribution 1, tuning variable initialization on a single MDP results in a training performance comparable to performance across the entire distribution. However, for distribution 2, the training performance of variable initialization on a single MDP is deceptive and constant initialization performs better on the full distribution. That is, variable initialization overfits. As more training samples for distribution 2 are provided, overfitting is mitigated and variable initialization is the better choice.

Since performance on the training set of MDP samples can be deceptive, we would like a way to bound the *generalization error*. The generalization error is the difference in the performance of a selected algorithm on a training set and its performance on the full distribution. More formally, let \mathcal{A} be the space of parameterized algorithms and let \mathcal{F} be the space of evaluation functions of each algorithm $a \in \mathcal{A}$ such that for all $f_a \in \mathcal{F}$, $f_a : \mathcal{D} \rightarrow [0, 1]$, where \mathcal{D} is the space of MDPs our MDP distribution spans. Let S be a sample z_1, \dots, z_m , chosen from a distribution D on \mathcal{D} . Then the generalization error of algorithm a is $|\hat{E}_S[f_a] - E[f_a]|$, where $\hat{E}_S[f_a] = \frac{1}{m} \sum_{i=1}^m f_a(z_i)$ is the expected performance of a on the training data, and $E[f_a] = \int_{\mathcal{D}} f_a(z) D(z) dz$ is the expected performance of a on the full distribution.

A major uniform convergence result for Rademacher complexity [7] states that with probability $1 - \delta$ over choices of S , for every function $f \in \mathcal{F}$, we have $E[f(z)] \leq \hat{E}_S[f] + 2R_S(\mathcal{F}) + 3\sqrt{\frac{\ln(2/\delta)}{2m}}$, where $R_S(\mathcal{F}) = E_{\bar{\sigma}}[\sup_{f \in \mathcal{F}} (\frac{1}{m} \sum_{i=1}^m f(z_i) \sigma_i)]$, is the empirical Rademacher complexity, and $\bar{\sigma} = (\sigma_1, \dots, \sigma_m)$ is a vector of m independent random variables, with $Pr(\sigma_i = 1) = Pr(\sigma_i = -1) = 1/2$. The role of the σ variables is to create a kind of output noise in the dataset.

Unfortunately, the sup operator requires exhaustive search over all possible learning algorithms, which is generally intractable. Our main result is a new generalization bound for when the sup operator is replaced with a tractable weak form




Figure 4: Generalization error for 5-state chain Distributions 1 and 2 using an ensemble of $l = 20$ learning algorithms, each chosen from a subset of $k = 750$ learning algorithms.

of optimization. Specifically, the weak optimization we use is to generate an ensemble of l learning algorithms. This ensemble is produced by selecting l subsets of A , each of size k and for each subset choosing the maximum performing algorithm. At test time, the agent would then randomly select one of the winning l algorithms for use.

Under this weak form optimization, we now have the following generalization error bound. With probability $1 - \delta$

$$E_{F_k}[E_z[f_S^k] - \hat{E}_S[f_S^k]] \leq \frac{2}{\ell} \sum_{j=1}^{\ell} R_S(F_k^j) + 5\sqrt{\frac{\ln(3/\delta)}{2m}},$$

where F_k^j is the set of k evaluation functions for the k algorithms of the j th subset used to form our ensemble. Due to space constraints, we have omitted the proof of this theorem.

Using this weak form of optimization, we can now compute generalization error bounds for our training data. Figure 4 shows how our estimated error bound in the 5-state chain problems using our weak optimizer tracks the true performance on the full distribution. It correctly predicts that on distribution 1 variable initialization should be preferred even with only one training sample, but on distribution 2, constant initialization should be preferred at first. The conservative nature of the bounds means that there is a need for a large training set before they can be convinced that variable initialization is not overfitting. Fortunately, the two sets of algorithm have similar performance.

References

- [1] Andrew G Barto, Steven J Bradtke, and Satinder P Singh. Learning to act using real-time dynamic programming. *Artificial Intelligence*, 72(1):81–138, 1995.
- [2] Anselm Blumer, Andrzej Ehrenfeucht, David Haussler, and Manfred K Warmuth. Learnability and the Vapnik-Chervonenkis dimension. *Journal of the ACM*, 36(4):929–965, 1989.
- [3] Mario Bollini, Stefanie Tellex, Tyler Thompson, Nicholas Roy, and Daniela Rus. Interpreting and executing recipes with a cooking robot. In *Proceedings of International Symposium on Experimental Robotics (ISER)*, 2012.
- [4] Adi Botea, Markus Enzenberger, Martin Müller, and Jonathan Schaeffer. Macro-ff: Improving ai planning with automatically learned macro-operators. *Journal of Artificial Intelligence Research*, 24:581–621, 2005.
- [5] Nicholas K. Jong. The utility of temporal abstraction in reinforcement learning. In *Proceedings of the Seventh International Joint Conference on Autonomous Agents and Multiagent Systems*, 2008.
- [6] Ross A. Knepper, Stefanie Tellex, Adrian Li, Nicholas Roy, and Daniela Rus. Single assembly robot in search of human partner: Versatile grounded language generation. In *Proceedings of the HRI 2013 Workshop on Collaborative Manipulation*, 2013.
- [7] Mehryar Mohri, Afshin Rostamizadeh, and Ameet Talwalkar. *Foundations of Machine Learning*. MIT press, 2012.
- [8] M Newton, John Levine, and Maria Fox. Genetically evolved macro-actions in ai planning problems. *Proceedings of the 24th UK Planning and Scheduling SIG*, pages 163–172, 2005.
- [9] Malcolm Stens. A Bayesian framework for reinforcement learning. In *Proceedings of the Seventeenth International Conference on Machine Learning*, 2000.
- [10] Richard S Sutton, Doina Precup, and Satinder Singh. Between mdps and semi-mdps: A framework for temporal abstraction in reinforcement learning. *Artificial intelligence*, 112(1):181–211, 1999.
- [11] R.S. Sutton and A.G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, Cambridge, MA, USA, 1 edition, 1998.